

Lokesh Punati,11514358

## Overfitting and Underfitting

[https://scikit-learn.org/stable/auto\\_examples/model\\_selection/plot\\_underfitting\\_overfitting.html](https://scikit-learn.org/stable/auto_examples/model_selection/plot_underfitting_overfitting.html)

### Program for understanding Overfitting and Underfitting

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

def true_fun(X):
    return np.cos(1.5 * np.pi * X)

np.random.seed(0)

n_samples = 30
degrees = [1, 4, 15]

X = np.sort(np.random.rand(n_samples))
y = true_fun(X) + np.random.randn(n_samples) * 0.1

plt.figure(figsize=(14, 5))
for i in range(len(degrees)):
    ax = plt.subplot(1, len(degrees), i + 1)
    plt.setp(ax, xticks=(), yticks=())


    polynomial_features = PolynomialFeatures(degree=degrees[i],
include_bias=False)
    linear_regression = LinearRegression()
    pipeline = Pipeline([
        ("polynomial_features", polynomial_features),
        ("linear_regression", linear_regression),
    ])
    pipeline.fit(X[:, np.newaxis], y)

    # Evaluate the models using crossvalidation
    scores = cross_val_score(
        pipeline, X[:, np.newaxis], y,
        scoring="neg_mean_squared_error", cv=10
    )
```

```

X_test = np.linspace(0, 1, 100)
plt.plot(X_test, pipeline.predict(X_test[:, np.newaxis]),
label="Model")
plt.plot(X_test, true_fun(X_test), label="True function")
plt.scatter(X, y, edgecolor="b", s=20, label="Samples")
plt.xlabel("x")
plt.ylabel("y")
plt.xlim((0, 1))
plt.ylim((-2, 2))
plt.legend(loc="best")
plt.title(
    "Degree {} \nMSE = {:.2e} (+/- {:.2e})".format(
        degrees[i], -scores.mean(), scores.std()
    )
)
plt.show()

```



## Overfitting (Printing accuracy at different steps)

<https://machinelearningmastery.com/overfitting-machine-learning-models/>

*# evaluate decision tree performance on train and test sets with different tree depths*

```

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
from matplotlib import pyplot

```

*# define dataset*

```

X, y = make_classification(n_samples=10000, n_features=20,
n_informative=5, n_redundant=15, random_state=1)

```

*# summarize the dataset*

```

print(X.shape, y.shape)

```

```

(10000, 20) (10000,)

```

```

# split into train test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3)
# summarize the shape of the train and test sets
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

(7000, 20) (3000, 20) (7000,) (3000,)

train_scores, test_scores = list(), list()

# define the tree depths to evaluate
values = [i for i in range(1, 31)]

# evaluate a decision tree for each depth
for i in values:
    # configure the model
    model = DecisionTreeClassifier(max_depth=i)
    # fit model on the training dataset
    model.fit(X_train, y_train)
    # evaluate on the train dataset
    train_yhat = model.predict(X_train)
    train_acc = accuracy_score(y_train, train_yhat)
    train_scores.append(train_acc)
    # evaluate on the test dataset
    test_yhat = model.predict(X_test)
    test_acc = accuracy_score(y_test, test_yhat)
    test_scores.append(test_acc)
    # summarize progress
    print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))

>1, train: 0.763, test: 0.767
>2, train: 0.804, test: 0.805
>3, train: 0.871, test: 0.868
>4, train: 0.906, test: 0.890
>5, train: 0.924, test: 0.901
>6, train: 0.937, test: 0.912
>7, train: 0.947, test: 0.917
>8, train: 0.956, test: 0.914
>9, train: 0.966, test: 0.917
>10, train: 0.975, test: 0.911
>11, train: 0.981, test: 0.913
>12, train: 0.985, test: 0.909
>13, train: 0.990, test: 0.909
>14, train: 0.993, test: 0.907
>15, train: 0.995, test: 0.905
>16, train: 0.996, test: 0.910
>17, train: 0.997, test: 0.908
>18, train: 0.998, test: 0.904
>19, train: 0.999, test: 0.905
>20, train: 0.999, test: 0.903
>21, train: 1.000, test: 0.902

```

```
>22, train: 1.000, test: 0.905
>23, train: 1.000, test: 0.903
>24, train: 1.000, test: 0.901
>25, train: 1.000, test: 0.901
>26, train: 1.000, test: 0.906
>27, train: 1.000, test: 0.905
>28, train: 1.000, test: 0.900
>29, train: 1.000, test: 0.900
>30, train: 1.000, test: 0.908
```

---

## Cross-validation

[https://scikit-learn.org/stable/modules/cross\\_validation.html](https://scikit-learn.org/stable/modules/cross_validation.html)

```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn import svm
```

```
X, y = datasets.load_iris(return_X_y=True)
X.shape, y.shape

((150, 4), (150,))
```

### Basic method to compute score

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.4, random_state=0)
```

```
X_train.shape, y_train.shape
```

```
X_test.shape, y_test.shape
```

```
clf = svm.SVC(kernel='linear', C=1).fit(X_train, y_train)
clf.score(X_test, y_test)
```

```
0.9666666666666667
```

**Estimate the accuracy by splitting the data, computing the score 5 consecutive times (with different splits each time)**

```
from sklearn.model_selection import cross_val_score
clf = svm.SVC(kernel='linear', C=1, random_state=42)
scores = cross_val_score(clf, X, y, cv=5)
scores
```

```
array([0.96666667, 1.          , 0.96666667, 0.96666667, 1.          ])

print("%0.2f accuracy with a standard deviation of %0.2f" %
      (scores.mean(), scores.std()))
```

0.98 accuracy with a standard deviation of 0.02

### Using the different scoring parameter

```
from sklearn import metrics
```

```
scores = cross_val_score(
    clf, X, y, cv=5, scoring='f1_macro')
scores
```

```
array([0.96658312, 1.          , 0.96658312, 0.96658312, 1.          ])
```

### Specified multiple metrics of predefined scorer names

```
from sklearn.model_selection import cross_validate
from sklearn.metrics import recall_score
```

```
scoring = ['precision_macro', 'recall_macro']
clf = svm.SVC(kernel='linear', C=1, random_state=0)
scores = cross_validate(clf, X, y, scoring=scoring)
sorted(scores.keys())
```

```
scores['test_recall_macro']
```

```
array([0.96666667, 1.          , 0.96666667, 0.96666667, 1.          ])
```

### Calculate cross validation score by passing a cross validation iterator

```
from sklearn.model_selection import ShuffleSplit
n_samples = X.shape[0]
cv = ShuffleSplit(n_splits=5, test_size=0.3, random_state=0)
cross_val_score(clf, X, y, cv=cv)
```

```
array([0.97777778, 0.97777778, 1.          , 0.95555556, 1.          ])
```

### Use an iterable yielding (train, test) splits as arrays of indices

```
def custom_cv_2folds(X):
    n = X.shape[0]
    i = 1
    while i <= 2:
        idx = np.arange(n * (i - 1) / 2, n * i / 2, dtype=int)
        yield idx, idx
        i += 1
```

```
custom_cv = custom_cv_2folds(X)
cross_val_score(clf, X, y, cv=custom_cv)
```

```
array([1.          , 0.97333333])
```

## Different type of Cross validation iterators

### K-fold

```
import numpy as np
from sklearn.model_selection import KFold
```

```
X = ["a", "b", "c", "d"]
kf = KFold(n_splits=2)
for train, test in kf.split(X):
    print("%s %s" % (train, test))
```

```
[2 3] [0 1]
[0 1] [2 3]
```

### Repeated K-Fold

```
import numpy as np
from sklearn.model_selection import RepeatedKFold
X = np.array([[1, 2], [3, 4], [1, 2], [3, 4]])
random_state = 12883823
rkf = RepeatedKFold(n_splits=2, n_repeats=2,
                    random_state=random_state)
for train, test in rkf.split(X):
    print("%s %s" % (train, test))
```

```
[2 3] [0 1]
[0 1] [2 3]
[0 2] [1 3]
[1 3] [0 2]
```

### Leave One Out (LOO)

```
from sklearn.model_selection import LeaveOneOut
```

```
X = [1, 2, 3, 4]
loo = LeaveOneOut()
for train, test in loo.split(X):
    print("%s %s" % (train, test))
```

```
[1 2 3] [0]
[0 2 3] [1]
[0 1 3] [2]
[0 1 2] [3]
```

### Leave P Out (LPO)

```
from sklearn.model_selection import LeavePOut
```

```
X = np.ones(4)
lpo = LeavePOut(p=2)
for train, test in lpo.split(X):
    print("%s %s" % (train, test))
```

```
[2 3] [0 1]
[1 3] [0 2]
[1 2] [0 3]
[0 3] [1 2]
[0 2] [1 3]
[0 1] [2 3]
```

Stratified k-fold

```
from sklearn.model_selection import StratifiedKFold, KFold
import numpy as np
```

```
X, y = np.ones((50, 1)), np.hstack(([0] * 45, [1] * 5))
skf = StratifiedKFold(n_splits=3)
```

```
for train, test in skf.split(X, y):
    print('train - {} | test - {}'.format(
        np.bincount(y[train]), np.bincount(y[test])))
```

```
train - [30  3] | test - [15  2]
train - [30  3] | test - [15  2]
train - [30  4] | test - [15  1]
```

```
kf = KFold(n_splits=3)
```

```
for train, test in kf.split(X, y):
    print('train - {} | test - {}'.format(
        np.bincount(y[train]), np.bincount(y[test])))
```

```
train - [28  5] | test - [17]
train - [28  5] | test - [17]
train - [34] | test - [11  5]
```

## GroupShuffleSplit

```
from sklearn.model_selection import GroupShuffleSplit
```

```
X = [0.1, 0.2, 2.2, 2.4, 2.3, 4.55, 5.8, 0.001]
```

```
y = ["a", "b", "b", "b", "c", "c", "c", "a"]
```

```
groups = [1, 1, 2, 2, 3, 3, 4, 4]
```

```
gss = GroupShuffleSplit(n_splits=4, test_size=0.5, random_state=0)
```

```
for train, test in gss.split(X, y, groups=groups):
    print("%s %s" % (train, test))
```

```
[0 1 2 3] [4 5 6 7]
[2 3 6 7] [0 1 4 5]
[2 3 4 5] [0 1 6 7]
[4 5 6 7] [0 1 2 3]
```

---

## Validation curve

[https://keeeto.github.io/blog/bias\\_variance/](https://keeeto.github.io/blog/bias_variance/)

[https://scikit-learn.org/stable/modules/learning\\_curve.html](https://scikit-learn.org/stable/modules/learning_curve.html)

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score, learning_curve,
validation_curve

df_train = pd.read_csv('train.csv')
df_test = pd.read_csv('test.csv')
df_comb = df_train.append(df_test)

X = pd.DataFrame()

def encode_sex(x):
    return 1 if x == 'female' else 0

def family_size(x):
    size = x.SibSp + x.Parch
    return 4 if size > 3 else size

X['Sex'] = df_comb.Sex.map(encode_sex)
X['Pclass'] = df_comb.Pclass
X['FamilySize'] = df_comb.apply(family_size, axis=1)

fare_median = df_train.groupby(['Sex', 'Pclass']).Fare.median()
fare_median.name = 'FareMedian'

age_mean = df_train.groupby(['Sex', 'Pclass']).Age.mean()
age_mean.name = 'AgeMean'

def join(df, stat):
    return pd.merge(df, stat.to_frame(), left_on=['Sex', 'Pclass'],
right_index=True, how='left')
```



```
X['Fare'] = df_comb.Fare.fillna(join(df_comb, fare_median).FareMedian)
X['Age'] = df_comb.Age.fillna(join(df_comb, age_mean).AgeMean)
```

```
def quantiles(series, num):
    return pd.qcut(series, num, retbins=True)[1]
```

```
def discretize(series, bins):
    return pd.cut(series, bins, labels=range(len(bins)-1),
include_lowest=True)
```

```
X['Fare'] = discretize(X.Fare, quantiles(df_comb.Fare, 10))
X['Age'] = discretize(X.Age, quantiles(df_comb.Age, 10))
```

```
X_train = X.iloc[:df_train.shape[0]]
X_test = X.iloc[df_train.shape[0]:]
```

```
y_train = df_train.Survived
```

```
clf_1 = RandomForestClassifier(n_estimators=100, bootstrap=True,
random_state=0)
clf_1.fit(X_train, y_train)
# Number of folds for cross validation
num_folds = 7
```

```
def plot_curve(ticks, train_scores, test_scores):
    train_scores_mean = -1 * np.mean(train_scores, axis=1)
    train_scores_std = -1 * np.std(train_scores, axis=1)
    test_scores_mean = -1 * np.mean(test_scores, axis=1)
    test_scores_std = -1 * np.std(test_scores, axis=1)

    plt.figure()
    plt.fill_between(ticks,
                     train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
color="b")
    plt.fill_between(ticks,
                     test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1,
color="r")
    plt.plot(ticks, train_scores_mean, 'b-', label='Training score')
    plt.plot(ticks, test_scores_mean, 'r-', label='Validation score')
    plt.legend(fancybox=True, facecolor='w')

    return plt.gca()
```

```
def plot_validation_curve(clf, X, y, param_name, param_range,
scoring='roc_auc'):
    plt.xkcd()
    ax = plot_curve(param_range, *validation_curve(clf, X, y,
cv=num_folds,
scoring=scoring,
```


```

param_name=param_name,

param_range=param_range, n_jobs=-1))
    ax.set_title('')
    ax.set_xticklabels([])
    ax.set_yticklabels([])
    ax.set_xlim(2,12)
    ax.set_ylim(-0.97, -0.83)
    ax.set_ylabel('Error')
    ax.set_xlabel('Model complexity')
    ax.text(9, -0.94, 'Overfitting', fontsize=22)
    ax.text(3, -0.94, 'Underfitting', fontsize=22)
    ax.axvline(7, ls='--')
    plt.tight_layout()

plot_validation_curve(clf_1, X_train, y_train, param_name='max_depth',
param_range=range(2,13))

```



---

## ROC

<https://www.statology.org/plot-roc-curve-python/>

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn import metrics
import matplotlib.pyplot as plt

#import dataset from CSV file on Github
url =
"https://raw.githubusercontent.com/Statology/Python-Guides/main/default.csv"
data = pd.read_csv(url)

#define the predictor variables and the response variable
X = data[['student', 'balance', 'income']]
y = data['default']

#split the dataset into training (70%) and testing (30%) sets
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.3,random_state=0)


#instantiate the model
log_regression = LogisticRegression()

#fit the model using the training data
log_regression.fit(X_train,y_train)

LogisticRegression()

#define metrics
y_pred_proba = log_regression.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)

#create ROC curve
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



**Task 1: Perform all of the above codes of Overfitting, Cross Validation, etc. with the help of the given reference link.**

**Task 2: Explain your analysis of the code. Make a detailed analysis that can also cover the following questions: (Submit the PDF of Report)**

- 1) According to you, why do overfitting and underfitting occur, and how resolve them? What is the difference between them?
- 2) What kind of pattern did you analyze in the Train and Test score while running the code of overfitting?
- 3) What is cross-validation, and what did you analyze in a different type of validation that you performed?
- 4) Explain the analysis from generated ROC and validation curve and what they represent?

**Task 3: Using the given Cross Validation iterators perform all types of Cross Validations we did in the task :**

- 1) K-fold
- 2) Repeated K-Fold
- 3) Leave One Out (LOO)

**Apart from this three, try to perform validation using three new iterators.**

**Task 4: With the help of the given code and references complete all of the following step:**

- 1) Choose one new dataset. Train a overfitted model with the help of any machine learning technique, such as KNN, classification, regression.
- 2) Try to resolve the overfitting.
- 3) Calculate the Validation score by any two or three given techniques and Validation iterators.
- 4) Generate the validation curve
- 5) Predict the output of testing data.

#### 6) Generate the ROC curve using the predicted data and actual data.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score

from sklearn.datasets import load_wine
X, y = make_classification(n_samples=5000, n_features=20,
n_informative=15)

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, shuffle=True)

print(len(X),len(X_train),len(X_test))

5000 4000 1000

from sklearn.neighbors import KNeighborsClassifier
# Create KNN classifier
knn = KNeighborsClassifier(n_neighbors = 3)
# Fit the classifier to the data
knn.fit(X_train,y_train)

KNeighborsClassifier(n_neighbors=3)

#show first 5 model predictions on the test data
knn.predict(X_test)[0:5]

array([0, 1, 1, 1, 1])

knn.score(X_test, y_test)

0.934

from sklearn.model_selection import ShuffleSplit
n_samples = X.shape[0]
cv = ShuffleSplit(n_splits=5, test_size=0.3, random_state=0)
cross_val_score(clf, X, y, cv=cv)

array([0.74066667, 0.73266667, 0.726        , 0.74866667, 0.72        ])

def custom_cv_2folds(X):
    n = X.shape[0]
    i = 1
    while i <= 2:
        idx = np.arange(n * (i - 1) / 2, n * i / 2, dtype=int)
        yield idx, idx
        i += 1
```

```

custom_cv = custom_cv_2folds(X)
cross_val_score(clf, X, y, cv=custom_cv)

array([0.748 , 0.7364])

import numpy as np
from sklearn.model_selection import KFold

X = ["a", "b", "c", "d"]
kf = KFold(n_splits=2)
for train, test in kf.split(X):
    print("%s %s" % (train, test))

[2 3] [0 1]
[0 1] [2 3]

import numpy as np
from sklearn.model_selection import RepeatedKFold
X = np.array([[1, 2], [3, 4], [1, 2], [3, 4]])
random_state = 12883823
rkf = RepeatedKFold(n_splits=2, n_repeats=2,
random_state=random_state)
for train, test in rkf.split(X):
    print("%s %s" % (train, test))

[2 3] [0 1]
[0 1] [2 3]
[0 2] [1 3]
[1 3] [0 2]

import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import load_wine
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import validation_curve

wine = load_wine()

# Create feature matrix and target vector
X, y = wine.data, wine.target

train_scores, test_scores = list(), list()

# define the tree depths to evaluate
values = [i for i in range(1, 31)]

for i in values:
    # configure the model
    model = RandomForestClassifier(max_depth=i)
    # fit model on the training dataset
    model.fit(X_train, y_train)

```

```

# evaluate on the train dataset
train_yhat = model.predict(X_train)
train_acc = accuracy_score(y_train, train_yhat)
train_scores.append(train_acc)
# evaluate on the test dataset
test_yhat = model.predict(X_test)
test_acc = accuracy_score(y_test, test_yhat)
test_scores.append(test_acc)
# summarize progress
print('>%d, train: %.3f, test: %.3f' % (i, train_acc, test_acc))
plt.show()

```

```

>1, train: 0.903, test: 0.893
>2, train: 0.922, test: 0.909
>3, train: 0.984, test: 0.960
>4, train: 0.996, test: 0.980
>5, train: 0.998, test: 0.987
>6, train: 0.999, test: 0.993
>7, train: 1.000, test: 0.993
>8, train: 1.000, test: 0.993
>9, train: 1.000, test: 0.996
>10, train: 1.000, test: 0.996
>11, train: 1.000, test: 0.993
>12, train: 1.000, test: 0.993
>13, train: 1.000, test: 0.996
>14, train: 1.000, test: 0.991
>15, train: 1.000, test: 0.991
>16, train: 1.000, test: 0.993
>17, train: 1.000, test: 0.991
>18, train: 1.000, test: 0.991
>19, train: 1.000, test: 0.993
>20, train: 1.000, test: 0.987
>21, train: 1.000, test: 0.993
>22, train: 1.000, test: 0.989
>23, train: 1.000, test: 0.996
>24, train: 1.000, test: 0.993
>25, train: 1.000, test: 0.996
>26, train: 1.000, test: 0.996
>27, train: 1.000, test: 0.993
>28, train: 1.000, test: 0.996
>29, train: 1.000, test: 0.993
>30, train: 1.000, test: 0.991

```

```

param_range = np.arange(1, 250, 2)

```

```

# Calculate accuracy on training and test set using range of parameter
values

```

```

train_scores, test_scores = validation_curve(RandomForestClassifier(),

```

X,



```

y,

param_name="n_estimators",

param_range=param_range,
cv=3,
scoring="accuracy",
n_jobs=-1)

# Calculate mean and standard deviation for training set scores
train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)

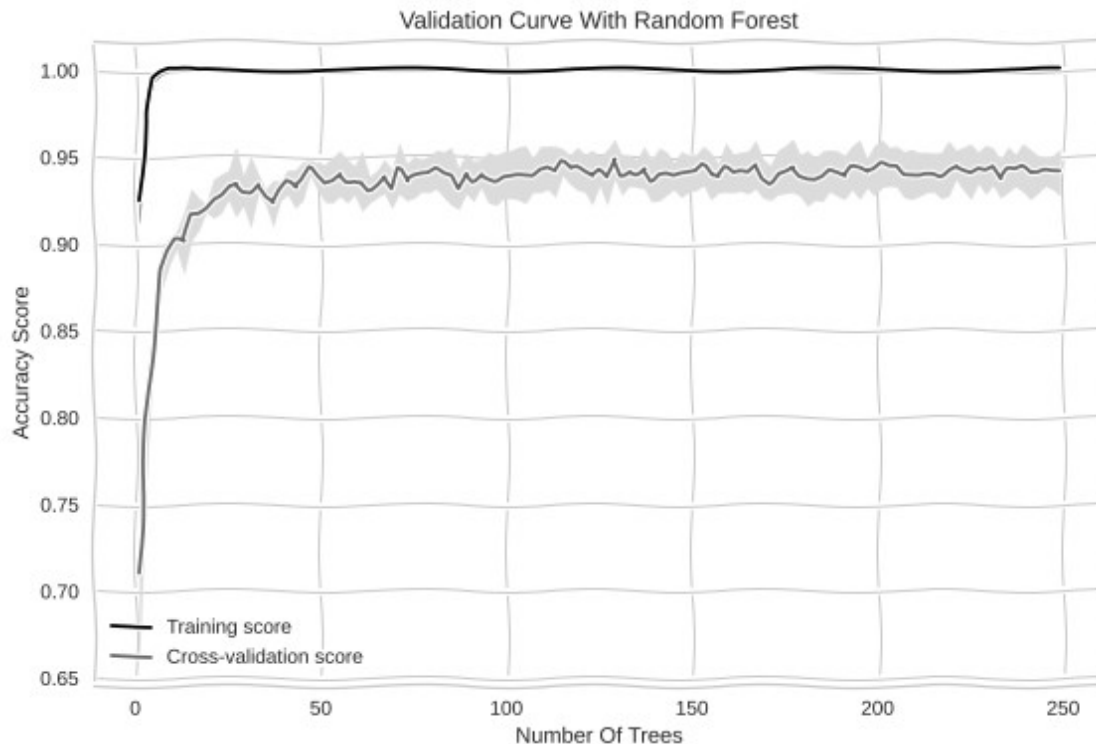
# Calculate mean and standard deviation for test set scores
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# Plot mean accuracy scores for training and test sets
plt.plot(param_range, train_mean, label="Training score",
color="black")
plt.plot(param_range, test_mean, label="Cross-validation score",
color="dimgrey")

# Plot accuracy bands for training and test sets
plt.fill_between(param_range, train_mean - train_std, train_mean +
train_std, color="gray")
plt.fill_between(param_range, test_mean - test_std, test_mean +
test_std, color="gainsboro")

# Create plot
plt.title("Validation Curve With Random Forest")
plt.xlabel("Number Of Trees")
plt.ylabel("Accuracy Score")
plt.tight_layout()
plt.legend(loc="best")
plt.show()

```



```
from sklearn.datasets import load_wine
wine = load_wine()
```

```
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(wine.data,
wine.target)
clf = GaussianNB()
clf.fit(X_train, y_train)
```

```
GaussianNB()
```

```
predicted = clf.predict(X_test)
expected = y_test
print(predicted)
```

```
[5 7 1 5 7 7 7 4 7 4 2 0 7 2 7 0 4 3 5 7 7 7 9 9 3 3 3 5 8 5 4 0 5 7 5
9 0
 8 7 2 2 1 6 1 4 6 7 9 6 2 7 1 6 8 7 7 7 7 3 8 2 8 1 8 4 2 0 5 7 7 3 1
8 1
 7 2 7 2 8 7 7 7 4 6 0 5 8 9 8 7 6 0 6 0 1 9 7 4 1 2 0 7 5 6 1 8 8 7 0
8 6
 0 5 6 4 3 5 6 5 0 6 3 7 1 1 3 4 7 6 6 9 8 4 6 6 4 3 2 7 6 6 6 6 3 8 7
8 6
 7 4 2 2 1 6 8 1 8 7 5 3 8 2 8 8 4 2 1 2 7 8 0 6 9 9 8 0 5 0 1 3 7 6 7
5 5]
```

```

7 7 6 0 6 2 2 1 7 4 9 7 9 1 8 8 7 6 8 2 1 7 8 2 2 6 8 2 3 7 9 4 1 4 3
9 6
8 8 9 0 6 5 9 9 6 6 7 5 6 5 2 8 7 8 7 8 3 7 5 7 8 0 9 4 1 0 3 6 6 0 2
7 4
8 6 0 7 4 9 1 7 7 2 4 8 9 0 8 6 9 2 1 9 2 1 0 2 0 4 5 9 3 7 3 6 3 8 1
6 5
1 0 2 8 5 4 6 6 2 1 9 6 1 4 8 3 9 4 5 5 2 4 7 3 3 6 1 2 9 0 9 5 5 0 1
1 6
5 8 4 9 2 8 2 0 5 4 0 7 4 7 3 2 8 0 6 7 6 5 4 8 5 7 6 0 5 2 3 2 1 7 2
5 7
1 7 6 4 4 7 0 1 0 1 8 6 5 5 5 5 2 6 7 6 3 2 7 7 3 9 5 4 1 6 8 3 2 2 0
7 2
8 1 7 4 8 9 2 0 3 1 3 4 8 5 4 8 3 8 1 8 5 7 9 6 3 9 9 9 8 1 3 3 5 0 2
0 4
4 0 8 7 1 3]

```

```

from sklearn import datasets

```

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

```

```

dataset = datasets.load_wine()
X = dataset.data
y = dataset.target

```

```

from sklearn import datasets, ensemble, metrics, model_selection,
dummy
import matplotlib.pyplot as plt

```

```

wine = datasets.load_wine()

```

```

# First we work only on sevens:
sevens = (wine.target == 7)

```

```

classifier = ensemble.RandomForestClassifier()
most_frequent = dummy.DummyClassifier(strategy='most_frequent')

```

```

X_train, X_test, y_train, y_test = model_selection.train_test_split(
    wine.data, sevens, random_state=0)

```

```

y_score = classifier.fit(X_train, y_train).predict_proba(X_test)
most_frequent_score = most_frequent.fit(X_train,
                                         y_train).predict_proba(X_test)

```

```

fpr, tpr, _ = metrics.roc_curve(y_test, y_score[:, 1])
roc_auc = metrics.auc(fpr, tpr)

```

```

fpr_dummy, tpr_dummy, _ = metrics.roc_curve(y_test,

```

```

most_frequent_score[:, 1])
roc_auc_dummy = metrics.auc(fpr_dummy, tpr_dummy)

plt.figure(figsize=(3, 3))
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='RandomForestClassifier\n(area = %0.2f)' %
         roc_auc)
plt.plot(fpr_dummy, tpr_dummy, color='.5',
         lw=lw, label='Dummy\n(area = %0.2f)' % roc_auc_dummy)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower left", title='ROC curve')
plt.tight_layout()
plt.show()

```

